

## Spark Installtion

```
# innstall java
!apt-get install openjdk-8-jdk-headless -qq > /dev/null

# install spark (change the version number if needed)
!wget -q https://archive.apache.org/dist/spark/spark-3.0.0/spark-3.0.0-bin-hadoop3.2.tgz

# unzip the spark file to the current folder
!tar xf spark-3.0.0-bin-hadoop3.2.tgz

# set your spark folder to your system path environment.
import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-3.0.0-bin-hadoop3.2"

# install findspark using pip
!pip install -q findspark
```

## Imports

```
import findspark
findspark.init()
import pyspark

from pyspark.sql import SparkSession
import pyspark.sql.functions as F
from pyspark.sql.types import *
from pyspark.sql.window import Window

from importlib import reload
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

## Spark Sesstion

```
def init_spark(app_name: str):
    spark = SparkSession.builder.appName(app_name).getOrCreate()
    sc = spark.sparkContext
    return spark, sc
```

```
spark, sc = init_spark('test')
```

```
spark, sc = init_spark( proj )
spark
```

## SparkSession - in-memory

### SparkContext

### [Spark UI](#)

Version

v3.0.0

Master

local[\*]

AppName

proj

## Place Based Analysis

```
username = "arieln"
password = "Qwerty12!"
server_name = "jdbc:sqlserver://technionddscourse.database.windows.net:1433"
database_name = "arieln"
url = server_name + ";" + "databaseName=" + database_name + ";"
df = spark.read \
    .format("com.microsoft.sqlserver.jdbc.spark") \
    .option("url", url) \
    .option("dbtable", "spatialDF") \
    .option("user", username) \
    .option("password", password).load()
df.show()
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount()

StationId	latitude	longitude	maxTemp	minTemp	avgPrecip
GME00121990	51.0892	12.9342	278.4731182795699	-41.03225806451613	16.22542989642368
UKE00105897	51.9589	1.0267	248.9777777777778	-9.466666666666667	17.05960073320977
SPE00120467	28.4775	-16.3292	290.9732142857143	94.9375	12.0054549621174
GME00132286	54.0206	9.9267	258.88659793814435	-34.329896907216494	19.28566346059702
GME00131386	52.3625	12.3883	284.5	-32.27450980392157	14.1052278387372
GME00102193	54.5289	9.5492	245.1078431372549	-16.833333333333332	25.7690426715061
GME00131794	52.0106	10.3956	273.3921568627451	-28.92156862745098	16.87486015946581
GME00123718	51.6267	10.3717	272.9306930693069	-41.99009900990099	18.39367412683995
GME00131938	48.7358	10.7428	265.6	-40.93684210526316	18.37077698957789
SPE00156522	42.1769	3.0967	304.6203703703704	8.74074074074074	13.81171679700619
FRM00007627	43.008	1.103	303.029702970297	-17.693069306930692	25.06517839477428
GME00128266	49.9203	8.9681	292.26	-34.45	14.84435285925015
NOE00109394	69.0589	18.5403	204.76415094339623	-198.41509433962264	18.07069425598378
GME00102356	49.7714	9.9589	287.5247524752475	-27.099009900990097	12.6146453613697
GME00122698	49.985	7.9544	288.2647058823529	-17.91176470588235	13.5107753418775
GME00129190	50.9789	12.3444	283.02912621359224	-40.699029126213595	14.8341867715139
UKM00003091	57.18	-2.2	190.8469387755102	-10.306122448979592	25.09809921029931
SPE00156171	41.3403	2.0803	315.35714285714283	43.526785714285715	16.74362695615368
SPE00119963	36.0153	-5.5975	272.80357142857144	89.49107142857143	13.31170862034153

```
|GME00122938| 49.665| 11.2253| 266.8910891089109| -45.37623762376238| 20.587414029596|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```




```
countries = ['GM','NO','SP','UK','FR']
df = df.withColumn("country", F.col("StationId").substr(1, 2))
countries_dict = {}
countries = {}
countries['GM'] = 'Germany'
countries['NO'] = 'Norway'
countries['SP'] = 'Spain'
countries['UK'] = 'United Kingdom'
countries['FR'] = 'France'


for country in countries:
    countries_dict[country] = df[df["country"] == country]
for country in countries:
    print("----- " + countries[country] + " -----")
    fig, axes = plt.subplots(nrows=1, ncols=3)
    sns.histplot(x="maxTemp", bins = 20, data = countries_dict[country].select('maxTemp').toPandas())
    axes[0].set_title("Max Temperature")
    axes[0].set_xlim(left=0, right=400)
    sns.histplot(x="minTemp", bins = 20, data = countries_dict[country].select('minTemp').toPandas())
    axes[1].set_title("Min Temperature")
    axes[1].set_xlim(left=-300, right=200)
    sns.histplot(x="avgPrcp", bins = 20, data = countries_dict[country].select('avgPrcp').toPandas())

    axes[2].set_title("Precipitation")
    axes[2].set_xlim(left=0, right=200)
    fig.tight_layout()
    plt.show()
```


## Germany




## Norway




## Spain

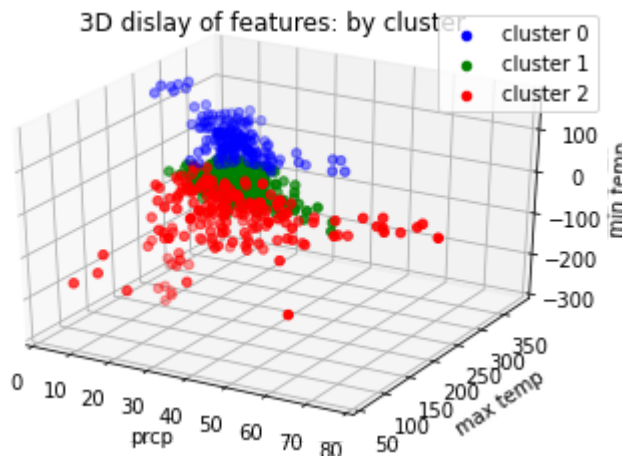


## United Kingdom






```
fig = plt.figure()
ax = fig.add_subplot(projection='3d')
colors = ["b", "g", "r", "m", "k"]
pd_df = df.toPandas()
for i in range(3):
    ax.scatter(pd_df[pd_df['cluster_label'] == i].avgPrcp, pd_df[pd_df['cluster_label'] == i].minTemp, c = colors[i], label = "cluster " + str(i))
ax.set_xlim(0,80)
ax.set_xlabel("prcp")
ax.set_ylabel("max temp")
ax.set_zlabel("min temp")
plt.legend()
plt.title("3D display of features: by cluster")
plt.show()
```



```
colors = [[0.4940,0.1840, 0.2560], [0.1940, 0.1440, 0.5060], [0.0940, 0.1240, 0.9960]]
reload(plt)
plt.xlim(0, 40)
plt.ylim(20, 80)
for i in range(3):
    pd_cdf = pd_df[pd_df['cluster_label']==i]
    sns.scatterplot(x = pd_cdf['longitude'],y = pd_cdf['latitude'],color = colors[i],label = "cluster " + str(i))
plt.legend()
```



## Time Based Analysis

```
from google.colab import drive
drive.mount('/content/gdrive')
df = spark.read \
    .format("com.microsoft.sqlserver.jdbc.spark") \
    .option("url", url) \
    .option("dbtable", "timeDF") \
    .option("user", username) \
    .option("password", password).load()
df.show()
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mou

```
+-----+-----+-----+
| StationId|      date|label|
+-----+-----+-----+
|SPE00156423|2016-01-03|    0|
|SPE00120125|2016-01-04|  408|
|SPE00120170|2016-01-05|   48|
|SPE00156585|2016-01-06|    0|
|SPE00119918|2016-01-10|    8|
|SPE00156369|2016-01-10|    0|
|SPE00156414|2016-01-10|   34|
|SPE00156648|2016-01-10|    0|
|SPE00120017|2016-01-11|   48|
|SPE00156018|2016-01-13|    0|
|SPE00156288|2016-01-13|    0|
|SPW00013024|2016-01-13|    0|
|SPE00156594|2016-01-15|    0|
|SPE00156513|2016-01-17|    0|
|SPE00156000|2016-01-18|    0|
|SPE00156189|2016-01-19|    0|
|SP000007038|2016-01-20|    0|
|SPE00156549|2016-01-22|    0|
|SPE00156162|2016-01-23|    0|
|SPE00156621|2016-01-24|    0|
+-----+-----+-----+
```

only showing top 20 rows

## Month Based

```

mount_agg_df = df.withColumn('month', F.split(df.date, '-')[1].cast('integer'))
mount_agg_df = mount_agg_df.withColumn('year', F.split(mount_agg_df['date'], '-')[0].cast('int'))
mount_agg_df = mount_agg_df.groupby('StationId', 'month', 'year').sum('label')
mount_agg_df = mount_agg_df.groupby('StationId', 'month').mean('sum(label)')
mount_agg_df = mount_agg_df.withColumn('PRCP', mount_agg_df['avg(sum(label))']).drop('avg(sum(label))')
mount_agg_df.show()

```

```

+-----+-----+-----+
| StationId|month| PRCP|
+-----+-----+-----+
|GME00115771| 8|1298.4|
|GME00124810| 8|1333.6|
|GME00121750| 2|1086.8|
|NOE00110554| 8|3561.2|
|SPE00156342| 4| 666.2|
|FR000007130|12|1446.8|
|NOE00134718|10| 841.0|
|GME00131530| 8| 911.2|
|NOE00109876| 1| 657.0|
|GME00125146| 1|1452.0|
|NO000001026|10|1630.0|
|GM000003038| 4| 653.6|
|GME00121654|10|1322.4|
|NOE00134574| 3| 400.4|
|NOE00133254| 7| 835.4|
|GME00132154| 3| 837.6|
|GME00123130| 2|1234.0|
|NOE00111004| 7| 809.6|
|GMM00010480| 1| 910.4|
|GME00129190| 6|1170.0|
+-----+-----+-----+
only showing top 20 rows

```

```

pd_df = mount_agg_df.toPandas()
sns.scatterplot(x = pd_df['month'], y = pd_df['PRCP'])
plt.show()

```



```

mount_agg_df = mount_agg_df.drop('StationId')
mount_agg_df = mount_agg_df.groupby("month").agg(F.min("PRCP"), F.max('PRCP'), F.var_pop("PRCP"))
mount_agg_df = mount_agg_df.withColumn('min_PRC', mount_agg_df['min(PRC)']).drop('min(PRC)')
mount_agg_df = mount_agg_df.withColumn('max_PRC', mount_agg_df['max(PRC)']).drop('max(PRC)')
mount_agg_df = mount_agg_df.withColumn('var_PRC', mount_agg_df['var_pop(PRC)']).drop('var_pop(PRC)')
mount_agg_df = mount_agg_df.withColumn('range_PRC', mount_agg_df.max_PRC - mount_agg_df.min_PRC)
mount_agg_df.show()


```

month	min_PRC	max_PRC	var_PRC	range_PRC
12	3.0	7406.8	823054.7028707761	7403.8
1	8.8	6458.0	595233.2888347387	6449.2
6	0.4	3980.8	341531.8432363311	3980.4
3	78.0	4798.8	323520.29335984966	4720.8
5	0.0	4702.4	316599.0529666013	4702.4
9	0.0	7847.2	663580.216593518	7847.2
4	22.0	3992.0	175490.68071661235	3970.0
8	0.0	6398.8	525633.3475020734	6398.8
7	0.0	4408.0	362539.1955357152	4408.0
10	0.0	6577.2	508345.22355317156	6577.2
11	0.0	5657.6	540652.4992342772	5657.6
2	57.2	4938.4	414570.52899540437	4881.2

```

pd_month_df = mount_agg_df.toPandas()
sns.scatterplot(x = pd_month_df['month'], y = pd_month_df['min_PRC'])
sns.scatterplot(x = pd_month_df['month'], y = pd_month_df['max_PRC'])
plt.show()

```




```


sns.scatterplot(x = pd_month_df['month'], y = pd_month_df['var_PRC'])
plt.show()

```





```
sns.scatterplot(x = pd_month_df['month'], y = pd_month_df['range_PRCP'])
plt.show()
```



```
sns.scatterplot(x = pd_month_df['range_PRCP'], y = pd_month_df['var_PRCP'])
x = pd_month_df['range_PRCP']
y = pd_month_df['var_PRCP']
z = np.polyfit(x, y, 1)
p = np.poly1d(z)
plt.plot(x, p(x), "r--")
plt.show()
```



## Time - Place Based Analysis

```
from google.colab import drive
drive.mount('/content/gdrive')
```

```
time_df = spark.read \
    .format("com.microsoft.sqlserver.jdbc.spark") \
    .option("url", url) \
    .option("dbtable", "timeDF") \
    .option("user", username) \
    .option("password", password).load()
time_df.show()
```

```
place_df = spark.read \
    .format("com.microsoft.sqlserver.jdbc.spark") \
    .option("url", url) \
    .option("dbtable", "spatialDF") \
    .option("user", username) \
    .option("password", password).load()
place_df.show()
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mou

```
+-----+-----+-----+
| StationId|      date|label|
+-----+-----+-----+
|SPE00156423|2016-01-03|    0|
|SPE00120125|2016-01-04|  408|
|SPE00120170|2016-01-05|   48|
|SPE00156585|2016-01-06|    0|
|SPE00119918|2016-01-10|    8|
|SPE00156369|2016-01-10|    0|
|SPE00156414|2016-01-10|   34|
|SPE00156648|2016-01-10|    0|
|SPE00120017|2016-01-11|   48|
|SPE00156018|2016-01-13|    0|
|SPE00156288|2016-01-13|    0|
|SPW00013024|2016-01-13|    0|
|SPE00156594|2016-01-15|    0|
|SPE00156513|2016-01-17|    0|
|SPE00156000|2016-01-18|    0|
|SPE00156189|2016-01-19|    0|
|SP000007038|2016-01-20|    0|
|SPE00156549|2016-01-22|    0|
|SPE00156162|2016-01-23|    0|
|SPE00156621|2016-01-24|    0|
```

+-----+-----+-----+

only showing top 20 rows

StationId	latitude	longitude	maxTemp	minTemp	avgPrecip
GME00121990	51.0892	12.9342	278.4731182795699	-41.03225806451613	16.22542989642368
UKE00105897	51.9589	1.0267	248.9777777777778	-9.466666666666667	17.05960073320977
SPE00120467	28.4775	-16.3292	290.9732142857143	94.9375	12.0054549621174
GME00132286	54.0206	9.9267	258.88659793814435	-34.329896907216494	19.28566346059702
GME00131386	52.3625	12.3883	284.5	-32.27450980392157	14.1052278387372
GME00102193	54.5289	9.5492	245.1078431372549	-16.833333333333332	25.7690426715061
GME00131794	52.0106	10.3956	273.3921568627451	-28.92156862745098	16.87486015946581
GME00123718	51.6267	10.3717	272.9306930693069	-41.99009900990099	18.39367412683995
GME00131938	48.7358	10.7428	265.6	-40.93684210526316	18.37077698957789
SPE00156522	42.1769	3.0967	304.6203703703704	8.74074074074074	13.81171679700619
FRM00007627	43.008	1.103	303.029702970297	-17.693069306930692	25.06517839477428
GME00128266	49.9203	8.9681	292.26	-34.45	14.84435285925015
NOE00109394	69.0589	18.5403	204.76415094339623	-198.41509433962264	18.07069425598378
GME00102356	49.7714	9.9589	287.5247524752475	-27.099009900990097	12.6146453613697
GME00122698	49.985	7.9544	288.2647058823529	-17.91176470588235	13.5107753418775
GME00129190	50.9789	12.3444	283.02912621359224	-40.699029126213595	14.8341867715139
UKM00003091	57.18	-2.2	190.8469387755102	-10.306122448979592	25.09809921029931
SPE00156171	41.3403	2.0803	315.35714285714283	43.526785714285715	16.74362695615368
SPE00119963	36.0153	-5.5975	272.80357142857144	89.49107142857143	13.31170862034153
GME00122938	49.665	11.2253	266.8910891089109	-45.37623762376238	20.5874140295963

+-----+-----+-----+

only showing top 20 rows

```
time_place_df = time_df.join(place_df.select('StationId', 'cluster_label'), on = 'StationId')
time_place_df.show()
```



StationId	date	label	cluster_label
SPE00156423	2016-01-03	0	0
SPE00120125	2016-01-04	408	1
SPE00120170	2016-01-05	48	0
SPE00156585	2016-01-06	0	0
SPE00119918	2016-01-10	8	0
SPE00156369	2016-01-10	0	0
SPE00156414	2016-01-10	34	0
SPE00156648	2016-01-10	0	0
SPE00120017	2016-01-11	48	0
SPE00156018	2016-01-13	0	2
SPE00156288	2016-01-13	0	0
SPW00013024	2016-01-13	0	0
SPE00156594	2016-01-15	0	0
SPE00156513	2016-01-17	0	1
SPE00156000	2016-01-18	0	0
SPE00156189	2016-01-19	0	0
SP000007038	2016-01-20	0	0
SPE00156549	2016-01-22	0	1
SPE00156162	2016-01-23	0	2
SPE00156621	2016-01-24	0	1

```
+-----+-----+-----+-----+
only showing top 20 rows
```

```
time_place_df = time_place_df.withColumn('date', F.substring(time_place_df.date, 0, 7))
time_place_df = time_place_df.withColumn('date', 10*F.split(time_place_df.date, '-')[0] + (10
time_place_df = time_place_df.groupby('StationId', 'date', 'cluster_label').sum('label')
time_place_df = time_place_df.withColumn('sum_PRCP', time_place_df['sum(label)']).drop('sum(1
time_place_df.show()
```

```
+-----+-----+-----+-----+
| StationId|          date|cluster_label|sum_PRCP|
+-----+-----+-----+-----+
|SPE00156018| 20162.30769230769|          2|    516|
|SPE00156351| 20160.76923076923|          1|    142|
|SPE00120089| 20162.30769230769|          0|    320|
|SPE00155964| 20161.53846153846|          0|    159|
|SPE00120035| 20163.846153846152|          0|    784|
|SPE00120467| 20162.30769230769|          0|   1296|
|SPE00120395| 20165.384615384617|          0|    128|
|SPE00120287| 20162.30769230769|          0|    780|
|SPE00156090| 20163.076923076922|          0|    598|
|SPE00156495| 20166.923076923078|          2|   1552|
|SPE00119855| 20164.615384615383|          0|     20|
|SPE00120062| 20168.46153846154|          0|   1644|
|SPE00120170| 20175.384615384617|          0|     20|
|SPE00120530| 20167.69230769231|          0|    732|
|SPE00119936| 20170.76923076923|          0|    248|
|SPM00008360| 20173.076923076922|          0|     46|
|SPE00155982| 20176.153846153848|          0|    186|
|SPE00156531| 20163.076923076922|          0|    706|
|SPE00173114| 20176.153846153848|          1|    483|
|SPE00155991| 20169.23076923077|          0|    333|
+-----+-----+-----+-----+
only showing top 20 rows
```

```
prcp_ts = time_place_df.toPandas()
for k in range(3):
    fig, ax1 = plt.subplots(1, figsize=(10,8))
    sns.lineplot(x="date", y="sum_PRCP", data = prcp_ts[prcp_ts['cluster_label']==k], ax=ax1)
    prcp_mean = prcp_ts[prcp_ts['cluster_label']==k]["sum_PRCP"].mean()
    prcp_std = prcp_ts[prcp_ts['cluster_label']==k]["sum_PRCP"].std()
    ax1.axhline(y = prcp_mean, color = 'r', linestyle = '-', label="prcp_mean")
    ax1.axhline(y = prcp_mean + 2*abs(prcp_std - prcp_mean), color = 'k', linestyle = 'dotted',
    ax1.axhline(y = prcp_mean - 2*abs(prcp_std - prcp_mean), color = 'k', linestyle = 'dotted')
    ax1.set_title("Percipitation from 2016-2021: cluster " + str(k))
```



```
labels = ['Cluster 0', 'Cluster 1', 'Cluster 2']
```

```
for k in range(3):
```

```
    sns.linelot(x="date", v="sum_PRCP", data = prcp_ts[prcp_ts['cluster_label']==k], label = 1
```

```
sum_prcp = data.groupby('date').sum_prcp  
plt.legend()  
plt.show()
```

